

Benchmark of error rate inference from UMI-tagged data and PCR error model

Mikhail Shugay

November 6, 2016

Load data from 2 independent experiments with 10 different PCR assays.

```
df <- data.frame()

for (q in c(20,25,30)){
  for (proj in c(73, 82)) {
    for (sample in c("encyclo", "kappa-hf-taq", "phusion", "sd-hs", "snp-detect",
                     "taq-hs", "tersus", "tersus-snp-buff", "truseq", "velox")) {

      .df <- read.table(
        paste(
          paste(q, "-1", paste("polerr", proj, sep=""), sep="_"),
          paste("polerr", proj, ".", sample, ".variant.caller.txt", sep=""), sep="/"),
        header=T, sep="\t", stringsAsFactors = F)

      .df$q <- q
      .df$proj <- proj
      .df$sample <- sample
      df <- rbind(df, .df)
    }
  }
}

df <- subset(df, freq > 0 & freq < 0.001 &
             !grepl("D", mutation) & !grepl("I", mutation) & global.est == 0 &
             coverage > 0)

df$mut.split <- sapply(df$mutation, function(x) strsplit(as.character(x), "[S:>]"))
df$mutation.pos <- as.integer(sapply(df$mut.split, function(x) x[2]))
df$mutation.from <- sapply(df$mut.split, function(x) x[3])
df$mutation.to <- sapply(df$mut.split, function(x) x[4])
df$mut.split <- NULL
```

Model estimates for error rate:

```
library(ggplot2)
```

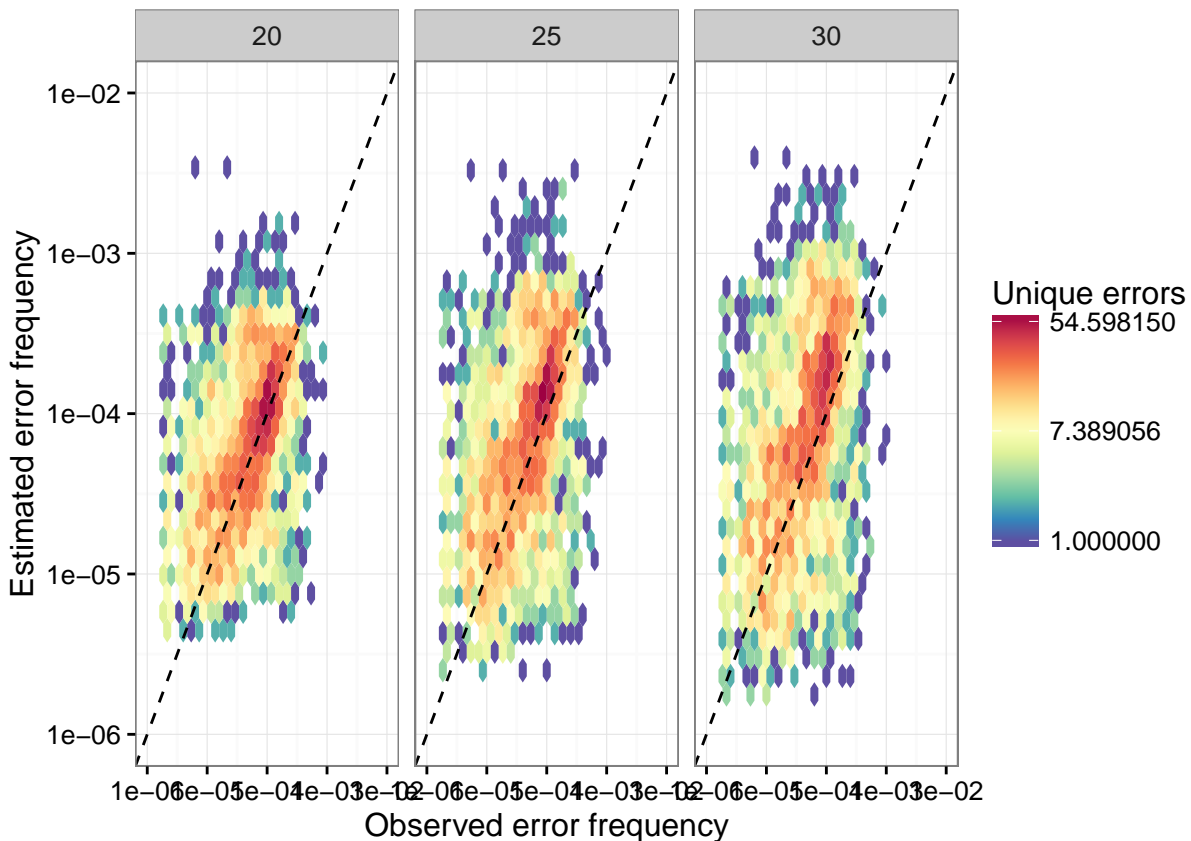
```
## Warning: package 'ggplot2' was built under R version 3.2.4
```

```
library(RColorBrewer)
rf <- colorRampPalette(rev(brewer.pal(11, 'Spectral')))(32)
r <- rf(32)

ggplot(df, aes(x=freq, y=error.rate)) +
  stat_binhex() +
```

```
geom_abline(intercept = 0, slope=1, linetype="dashed") +
scale_x_log10("Observed error frequency", limits=c(1e-6, 1e-2),
              breaks=c(1e-6,1e-5,1e-4,1e-3,1e-2)) +
scale_y_log10("Estimated error frequency", limits=c(1e-6, 1e-2),
              breaks=c(1e-6,1e-5,1e-4,1e-3,1e-2)) +
facet_wrap(~q) +
scale_fill_gradientn("Unique errors", colors=r, trans="log") +
theme_bw()
```

```
## Warning: package 'hexbin' was built under R version 3.2.5
```



```
for (qq in c(20, 25, 30)) {
  print(with(subset(df, q == qq), cor(freq, error.rate, method = "spearman")))
}
```

```
## [1] 0.4343872
## [1] 0.4344433
## [1] 0.4356582
```

We'll use Bayesian model when computing P-values (and Q-scores) of erroneous variants based on estimated error rate and their count, accounting for a high degree of uncertainty in error rate estimates. Thus, we'll compute P-values based on Negative Binomial distribution instead of Poisson model. To estimate its parameters, let's fit Gamma distribution (conjugate prior of Poisson distribution) to $P(\lambda | \text{error rate estimate} * \text{coverage})$. First let's plot the conditional probability:

```
library(plyr)

df.1 <- subset(df, q == 20)

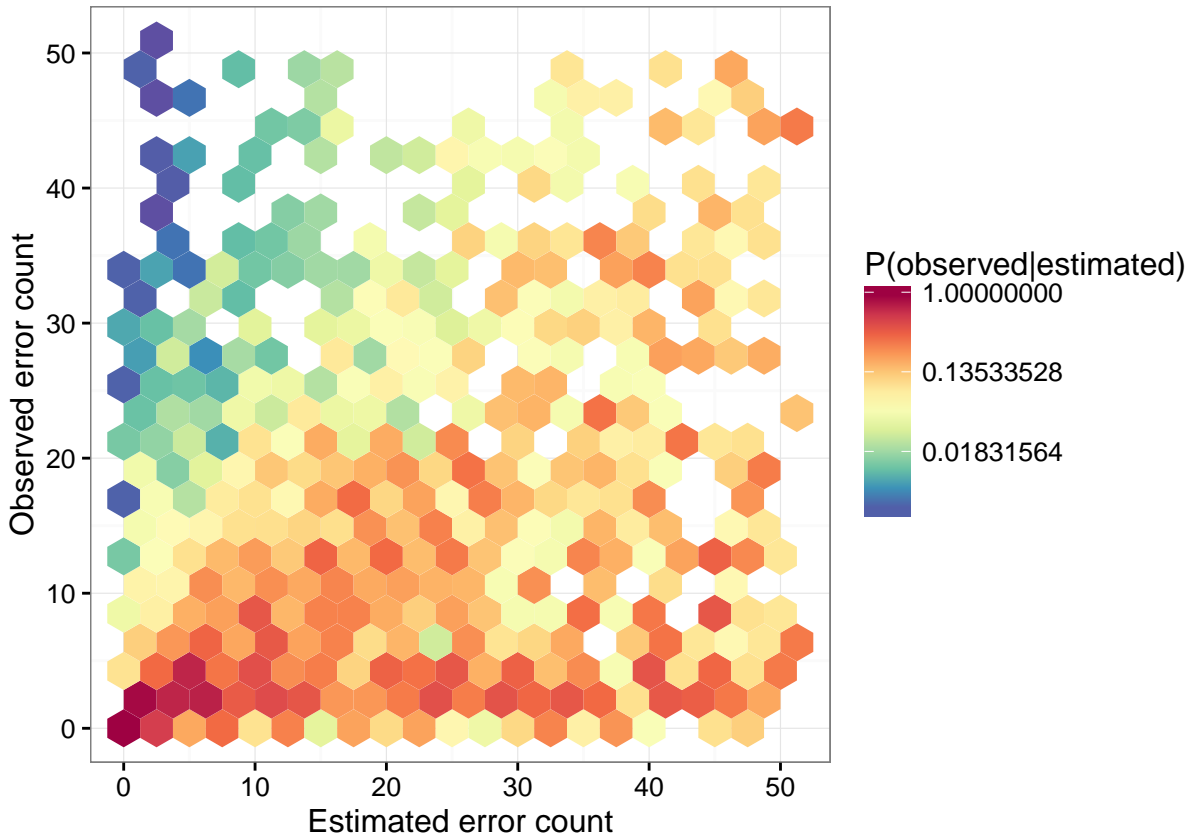
df.1$count.est <- round(df.1$coverage * df.1$error.rate)
df.1$count.obs <- round(df.1$coverage * df.1$freq)
df.1 <- subset(df.1, count.est <= 50 & count.obs <= 50)

# Compute mean and variance for lambda, we'll need it later
df.1 <- ddply(df.1, .(count.est), transform, f.mean = mean(freq * coverage))
df.1 <- ddply(df.1, .(count.est), transform, f.var = var(freq * coverage))

df.1.s <- ddply(df.1, .(count.est, count.obs),
  summarize,
  weight = length(count.est), f.mean = f.mean[1], f.var = f.var[1])
df.1.m <- ddply(df.1.s, .(count.est), transform, weight.norm = weight / sum(weight))

# requires install.packages("hexbin")


ggplot(df.1.m, aes(x=count.est, y=count.obs, weight=weight.norm)) +
  geom_hex(bins = 20) +
  xlab("Estimated error count") +
  ylab("Observed error count") +
  scale_fill_gradientn("P(observed|estimated)", colors=r, trans="log") +
  theme_bw()
```



Gamma distribution parameters α and β can be estimated through mean and variance of observed count of erroneous variants (λ). We'll first fit estimated error count to these parameters:

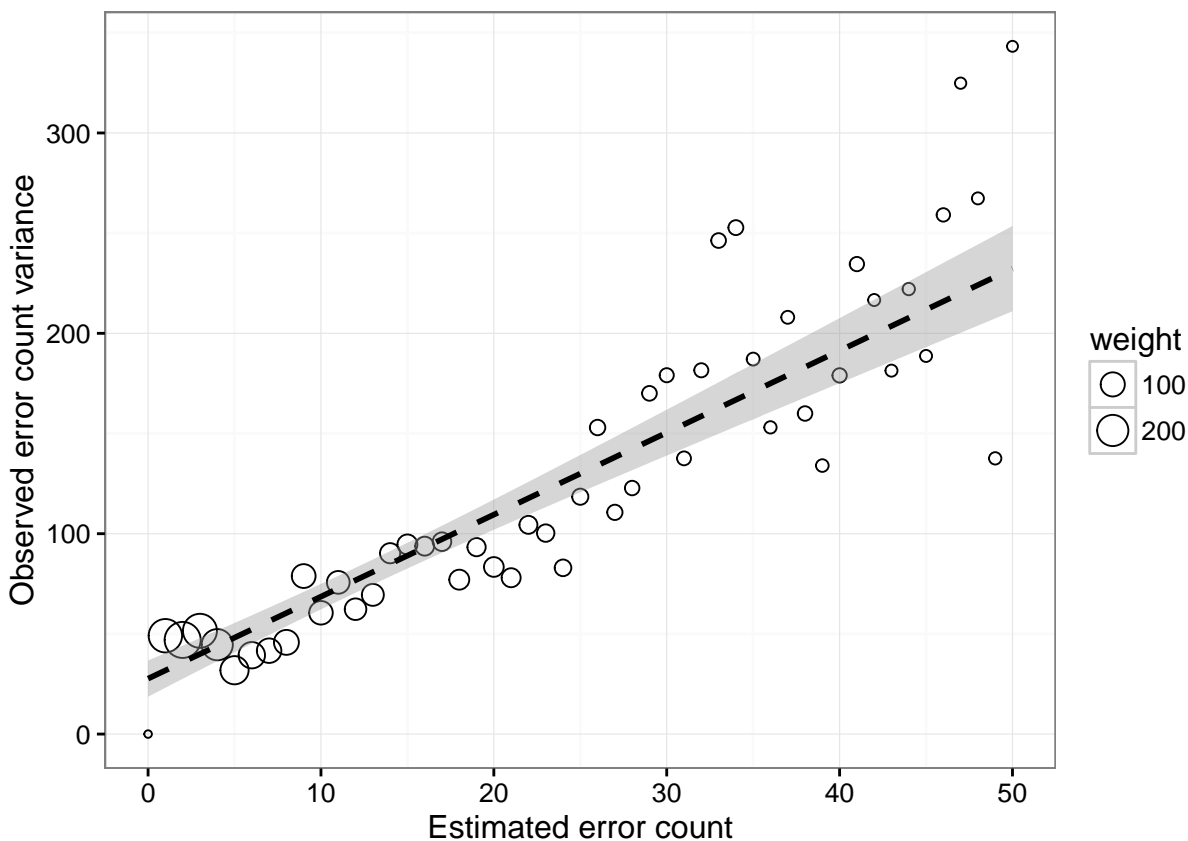
```
df.1.m.0 <- ddply(df.1.m, .(count.est, f.mean, f.var),
  summarize,
  weight = sum(weight))

ggplot(df.1.m.0, aes(x=count.est, y=f.mean)) + geom_point(aes(size=weight), shape=21) +
  geom_smooth(aes(weight=weight), method=lm, color="black", linetype="dashed") +
  xlab("Estimated error count") +
  ylab("Observed error count mean") +
  theme_bw()
```



```
l.mean.fit <- lm(f.mean ~ count.est,
  df.1.m.0, weights=weight)

ggplot(df.1.m.0, aes(x=count.est, y=f.var)) + geom_point(aes(size=weight), shape=21) +
  geom_smooth(aes(weight=weight), method=lm, color="black", linetype="dashed") +
  xlab("Estimated error count") +
  ylab("Observed error count variance") +
  theme_bw()
```



```
l.var.fit <- lm(f.var ~ count.est,
               df.1.m.0, weights=weight)
```

Compute probabilities based on Gamma distribution, compare fitted and observed $P(\lambda | \text{error rate estimate} * \text{coverage})$:

```
df.1.m <- ddply(df.1.m, .(count.est),
               transform,
               f.mean.est =
                 coef(l.mean.fit)["(Intercept)"]
                 + coef(l.mean.fit)["count.est"] * count.est,
               f.var.est =
                 coef(l.var.fit)["(Intercept)"]
                 + coef(l.var.fit)["count.est"] * count.est)


df.1.m$a <- df.1.m$f.mean.est * df.1.m$f.mean.est / df.1.m$f.var.est
df.1.m$b <- df.1.m$f.mean.est / df.1.m$f.var.est
df.1.m$prob.gamma <- dgamma(df.1.m$count.obs, rate = df.1.m$b, shape = df.1.m$a)

library(reshape2)

df.1.m.1 <- melt(df.1.m, id.vars = c("count.est", "count.obs"),
                 measure.vars = c("prob.gamma", "weight.norm"))

# requires install.packages("quantreg")
```

```
ggplot(df.1.m.1, aes(x=count.est, group = interaction(variable,count.est), fill=variable)) +
  geom_boxplot(aes(y=count.obs, weight=value)) +
  xlab("Estimated error count") +
  ylab("Observed error count") +
  scale_fill_brewer(palette = "Set1", "Density") +
  theme_bw()
```



Lets now compute P-values for errors using fitted Negative Binomial model and compare them to real P-values of error counts:

```
df.1.m$size <- df.1.m$a + df.1.m$count.obs
df.1.m$prob <- 1/(1+df.1.m$b)

df.1.m$pval <- with(df.1.m,
  1.0 - pbinom(count.obs,
    prob = prob,
    size = round(size)) +
  0.5 * dbinom(count.obs, prob = prob, size = round(size)))

df.1.m$pval.true <- sapply(df.1.m$count.obs,
  function(x) sum(ifelse(x > df.1.m$count.obs, 0,
    ifelse(x == df.1.m$count.obs,
      0.5 * df.1.m$weight,
      df.1.m$weight)
    ))) / sum(df.1.m$weight)
```

```
ggplot(df.1.m, aes(x=-10*log10(pval), y=-10*log10(pval.true))) +
  stat_density2d(aes(fill=..level.., weight=weight), geom="polygon") +
  geom_abline(intercept = 0, slope = 1, color = "black", linetype="dashed") +
  scale_x_continuous("Computed Q score", limits=c(0,20)) +
  scale_y_continuous("True Q score", limits=c(0,20)) +
  scale_fill_gradientn(colors=r) +
  theme_bw()
```

